

# Lectures 5 & 6: Linear Regression

Isabel Valera

Machine Learning Group  
Department of Mathematics and Computer Science  
Saarland University, Saarbrücken, Germany

03.05.2021 & 05.05.2021

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 Appendix

# Main references

- Bishop - Chapter 3
- ESL - Chapter 3

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 Appendix

# Regression setting


**Regression setting:** We consider problems where the output (target) variable is real-valued, i.e.,  $\mathcal{Y} = \mathbb{R}$ . Moreover, we assume we have access to training data  $(X_i, Y_i)_{i=1}^n$ , which is an i.i.d. sample from the probability measure  $P$  on  $\mathcal{X} \times \mathcal{Y}$ . Note that here we treat each observation  $(X_i, Y_i)$  in the training dataset as a random variable.

**Goal:** Learn a mapping function  $f^*(X)$  that minimizes the risk  $R(f) = \mathbb{E}[L(Y, f(X))]$  with  $f(X) \in \mathcal{F}$ .

**Standard loss** is the squared loss, i.e.,  $L(y, f(\mathbf{x})) = (y - f(\mathbf{x}))^2$ , leading to a **least squares regression** problem.

**Linear regression** considers a family of regression functions  $\mathcal{F}$  that are linear, i.e., it takes the form  $\{\langle \mathbf{w}, \mathbf{x} \rangle + b, \text{ with } \mathbf{x}, \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R}\}$ .

# Illustration of Linear LSR



**FIGURE 3.1.** *Linear least squares fitting with  $X \in \mathbb{R}^2$ . We seek the linear function of  $X$  that minimizes the sum of squared residuals from  $Y$ .*

Figure: Figure from ESL book.

# Least squares regression

Least squares regression (LSR) considers the **Risk of squared loss**, i.e.,

$$\mathbb{E}[(Y - f(X))^2] = \mathbb{E}[\mathbb{E}[(Y - f(X))^2 | X]].$$

The optimal solution to LSR, i.e., the **Bayes optimal function**, is given by:

$$f^*(\mathbf{x}) = \mathbb{E}[Y | X = \mathbf{x}].$$

## Definition (Least Squares Regression)

Given a training sample  $D_n = (\mathbf{x}_i, y_i)_{i=1}^n$  with  $\mathbf{x}_i \in \mathcal{X}$  and  $y_i \in \mathbb{R}$ , and a function space  $\mathcal{F}$  we define **least squares regression** solution as

$$f_n^* = \arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n (y_i - f(\mathbf{x}_i))^2.$$

# Linear LSR

**Linear least squares regression** assumes a linear function class, i.e.,

$$\mathcal{F} = \left\{ f \mid f(x) = \sum_{i=1}^d w_i x_i + b = \langle \mathbf{w}, \mathbf{x} \rangle + b, \quad \mathbf{w} \in \mathbb{R}^d, b \in \mathbb{R} \right\}.$$

**Notation:**

- summarize the outputs  $(y_i)_{i=1}^n$  into a column vector  $\mathbf{Y} \in \mathbb{R}^n$  and the input vectors  $(\mathbf{x}_i)_{i=1}^n$  into a feature matrix  $\mathbf{X} \in \mathbb{R}^{n \times d}$ ,

$$\mathbf{Y} = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix}, \quad \mathbf{X} = \begin{pmatrix} x_{11} & \dots & x_{1d} \\ \vdots & & \vdots \\ x_{n1} & \dots & x_{nd} \end{pmatrix}.$$

- $\mathbf{w}$  is the **weight vector** and  $b$  is the **bias term** (a.k.a. intercept).

*Note:* During the course, we consider vectors as column vectors.

*Observation:* One may alternatively include the bias term in the weight vector by adding an additional column with all values equal to one to the feature matrix.





# Solution of Linear LSR

## Proposition

Let  $\mathbf{X} \in \mathbb{R}^{n \times d}$ . The solution  $\mathbf{w}_n$  of linear least squares regression is given by

$$\mathbf{w}_n = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y},$$

where the inverse  $(\mathbf{X}^T \mathbf{X})^{-1}$  exists if  $\mathbf{X}$  has rank  $d$ .

If  $\mathbf{X}$  has rank lower than  $d$ , then  $(\mathbf{X}^T \mathbf{X})^{-1}$  has to be understood in the sense of a generalized inverse. In this case the solution is not unique but if  $\mathbf{w}_n^1, \mathbf{w}_n^2$  are two different solutions, then their predictions agree on the training data

$$f_{\mathbf{w}_n^1}(\mathbf{x}_i) = \langle \mathbf{w}_n^1, \mathbf{x}_i \rangle = \langle \mathbf{w}_n^2, \mathbf{x}_i \rangle = f_{\mathbf{w}_n^2}(\mathbf{x}_i), \quad \forall i = 1, \dots, n.$$

Note: See 9 Appendix for details on the generalized inverse.

# Linear LSR solution Proof

**Proof:** Objective function of the optimization problem with  $\mathbf{w} \in \mathbb{R}^d$ ,

$$R_{LLSR}(\mathbf{w}) = \frac{1}{n} \sum_{i=1}^n (y_i - \langle \mathbf{w}, \mathbf{x}_i \rangle)^2 = \frac{1}{n} \|\mathbf{Y} - \mathbf{X}\mathbf{w}\|^2.$$

Taking the derivative with respect to  $\mathbf{w}$ ,

$$\nabla_{\mathbf{w}} R_{LLSR} = -\frac{2}{n} \mathbf{X}^T (\mathbf{Y} - \mathbf{X}\mathbf{w}).$$

The necessary condition for an extremum of  $R_{LLSR}$  is therefore

$$\frac{2}{n} \mathbf{X}^T (\mathbf{Y} - \mathbf{X}\mathbf{w}) = 0 \quad \rightarrow \quad \mathbf{X}^T \mathbf{Y} = (\mathbf{X}^T \mathbf{X}) \mathbf{w} \rightarrow \mathbf{w}_n = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}.$$

The solution  $\mathbf{w}_n$  is unique, if the Hessian of  $\mathbf{X}^T \mathbf{X}$  is positive-definite, which occurs if  $\mathbf{X}$  has rank  $d$ . If  $\mathbf{X}$  has rank smaller than  $d$ , then the solution is not unique.

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions**
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 Appendix

# Basis functions

**Basis functions/Feature maps:** Can be used to perform linear regression when the dependence between the features and outputs is non-linear but we have prior knowledge on the “non-linearity”.

- The idea is to map the input  $x \rightarrow \phi(x)$  into a new space in which the relationship between the new features  $X'$  and the outcome  $Y$  is linear.
- Examples:
  - $\mathcal{X} = \mathbb{R}$ , then  $\phi(x) = 1, x, x^2, x^3, \dots$  (polynomials),
  - $\mathcal{X} = [0, 2\pi]$ , then  $\phi(x) = \sin(x), \cos(x), \sin(2x), \cos(2x), \dots$  (Fourier basis).

Assuming a (fixed) pre-defined set of  $M$  **basis functions**,  $\phi_i : \mathbb{R}^d \rightarrow \mathbb{R}$ , we can define the function space as:

$$\mathcal{F} = \left\{ f : \mathbb{R}^M \rightarrow \mathbb{R}, f(\mathbf{x}) = \sum_{i=1}^M w_i \phi_i(\mathbf{x}) \mid \mathbf{w} \in \mathbb{R}^M \right\},$$

which remains linear with respect to the weight vector  $\mathbf{w}$ .

# LSR (generalization)

**Generalized feature matrix:**  $\Phi \in \mathbb{R}^{n \times M}$ ,

$$\Phi = \begin{pmatrix} \phi_1(X_1) & \dots & \phi_M(X_1) \\ \vdots & & \vdots \\ \phi_1(X_n) & \dots & \phi_M(X_n) \end{pmatrix},$$

**Least squares regression problem:**

$$\mathbf{w}_n = \arg \min_{\mathbf{w} \in \mathbb{R}^M} \frac{1}{n} \sum_{i=1}^n (Y_i - \langle \mathbf{w}, \Phi(X_i) \rangle)^2 = \frac{1}{n} \|\mathbf{Y} - \Phi \mathbf{w}\|^2,$$

with **solution**:

$$\mathbf{w}_n = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{Y},$$

where the matrix  $(\Phi^T \Phi)^{-1} \Phi^T \in \mathbb{R}^{M \times n}$  is the pseudo-inverse of  $\Phi$  (see 9 Appendix).


# The problem remains linear in $w$ !

## Some interesting properties:

- The final function,  $f(\mathbf{x}) = \langle \mathbf{w}_n, \Phi(\mathbf{x}) \rangle = \sum_{i=1}^M w_{ni} \phi_i(\mathbf{x})$ , is linear in the parameter  $\mathbf{w}_n$ . The functions  $\Phi(\mathbf{x})$  allow us to directly incorporate our prior knowledge.
- *Problem:* if we aim to model all polynomials in  $\mathbb{R}^d$ , then we need  $d$  polynomials of degree one (linear functions),  $\frac{d(d+1)}{2}$  polynomials of degree two, .... That is, the set of basis functions increases rapidly with  $d$ , making this approach unpractical.

# Example

Gas mileage dataset: input  $x$  is measures the horsepower, output measures the miles per gallon (mpg),  $n = 397$ .



$$\text{mpg} = \beta_0 + \beta_1 \times \text{horsepower} + \beta_2 \times \text{horsepower}^2 + \epsilon$$

Figure: Figure from EML course (Prof. Vreeken & Valera).



# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 **Bias-variance**
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 Appendix

# Relation to Bayes optimal function

Solutions  $\mathbf{w}_n$  of least squares are estimators for the optimal parameter  $\mathbf{w}^*$  (Bayes optimal **linear** function for the squared loss), i.e.,

$$\mathbf{w}^* = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \mathbb{E}[(Y - \langle \mathbf{w}, X \rangle)^2] = \mathbb{E}\left[\left(Y - \sum_{i=1}^d w_i X_i\right)^2\right],$$

whose optimal solution is given by (note that here  $X$  is a  $d$ -dimensional column vector random variable):

$$\mathbf{w}^* = \left(\mathbb{E}[X^T X]\right)^{-1} \mathbb{E}[X^T Y].$$

The empirical solutions  $\mathbf{w}_n$  depend on the training sample  $D = (X_i, Y_i)_{i=1}^n$ . **Key questions:**

- Is the average estimator  $\mathbf{w}_n$  over training samples of size  $n$  equal to the optimal  $\mathbf{w}^*$ ?
- How much does the estimator  $\mathbf{w}_n$  fluctuate around its average value over all possible training samples of size  $n$ ?

# Bias and Variance of an estimator

## Definition

Given a sample  $D = (X_i)_{i=1}^n$  and an estimator (also called statistics)  $\hat{y}_n : D \rightarrow \mathbb{R}$  of a quantity  $y \in \mathbb{R}$ . Then, the **bias** of the estimator  $\hat{y}_n$  is defined as

$$\text{Bias}(\hat{y}_n) = \mathbb{E}_D[\hat{y}_n] - y,$$

which corresponds to the difference of the expectation of  $\hat{y}_n$  over all training sets  $D$  (all possible i.i.d. training sets of size  $n$ ) and the true quantity  $y$ .

- The estimator  $\hat{y}_n$  is said to be **unbiased** if the bias is zero.
- It is **asymptotically unbiased** if  $\lim_{n \rightarrow \infty} \text{Bias}(\hat{y}_n) = 0$ .

The **variance** of an  $f_n$  is defined as,

$$\text{Var } \hat{y}_n = \mathbb{E}_D[(\hat{y}_n - \mathbb{E}_D[\hat{y}_n])^2],$$

and accounts for the variability of the estimate across different datasets of size  $n$ .

# Examples for bias and variance

- **The empirical mean**  $\mathbb{E}_{P_n}[X] = \frac{1}{n} \sum_{i=1}^n X_i$  is an **estimator of the true mean**  $\mathbb{E}_P[X] = \mathbb{E}[X]$ , and it fulfills that:

$$\mathbb{E}_D \left[ \frac{1}{n} \sum_{i=1}^n X_i \right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}_{X_i}[X_i] = \frac{1}{n} n \mathbb{E}[X] = \mathbb{E}[X] \implies \text{unbiased!}$$

- **The empirical variance**  $\text{Var}_{P_n}[X] = \frac{1}{n} \sum_{i=1}^n (X_i - \mathbb{E}_{P_n}[X])^2$  as an estimator of the **true variance**  $\text{Var}_P[X] = \text{Var}[X]$ , and it fulfills that:

$$\mathbb{E}_D [\text{Var}_{P_n}[X]] = \frac{n-1}{n} \text{Var}[X] \implies \text{biased! underestimation!}$$

Instead, the estimator  $\frac{1}{n-1} \sum_{i=1}^n (X_i - \mathbb{E}_{P_n}[X])^2$  for the variance of  $X$  is unbiased.

# Risk of an estimator

Let's assume a (fixed) estimator  $f_n$  of a regression function  $f : \mathcal{X} \rightarrow \mathcal{Y}$  and the squared loss, so that the Bayes optimal function is  $f^*(X) = \mathbb{E}[Y|X]$ . Then, the risk (i.e., the expected squared loss)  $R(f_n)$  of the estimator  $f_n$  is given by:

$$\begin{aligned} R(f_n) &= \mathbb{E}[(Y - f_n(X))^2] = \mathbb{E}[\mathbb{E}[(Y - f_n(X))^2|X]] \\ &= \mathbb{E}[\mathbb{E}[(Y - \mathbb{E}[Y|X] + \mathbb{E}[Y|X] - f_n(X))^2|X]] \\ &= \mathbb{E}[\mathbb{E}[(Y - \mathbb{E}[Y|X])^2|X]] + \mathbb{E}[\mathbb{E}[(\mathbb{E}[Y|X] - f_n(X))^2|X]] \\ &\quad + 2\mathbb{E}[\mathbb{E}[(Y - \mathbb{E}[Y|X])(\mathbb{E}[Y|X] - f_n(X))|X]], \\ &= \mathbb{E}[(Y - \mathbb{E}[Y|X])^2] + \mathbb{E}[(\mathbb{E}[Y|X] - f_n(X))^2] \end{aligned}$$

## Interpretation:

- The first term is the **Bayes optimal risk** (often also called noise term), where  $f^*(x) = \mathbb{E}[Y|X = x]$  is the Bayes optimal function for the squared loss (note that this term cannot be further reduced).
- The second term measures the **deviation of  $f_n$  from the Bayes optimal function**. It is a random quantity since  $f_n$  depends on the training data!

# Expected risk over all training datasets

**Expected risk**  $\mathbb{E}_D[R(f_n)]$  over all possible training sets  $D$  of size  $n$  is given by:

$$\mathbb{E}_D[R(f_n)] = \mathbb{E}[(Y - f^*(X))^2] + \mathbb{E}_D[\mathbb{E}_X[(f^*(X) - f_n(X))^2]],$$

Note that the first term is constant and the second (by exchanging the two expectations) can be elaborated as:

$$\begin{aligned} & \mathbb{E}_X[\mathbb{E}_D[(f_n(X) - f^*(X))^2]] \\ &= \mathbb{E}_X[\mathbb{E}_D[(f_n(X) - \mathbb{E}_D[f_n(X)] + \mathbb{E}_D[f_n(X)] - f^*(X))^2]] \\ &= \mathbb{E}_X[\mathbb{E}_D[(f_n(X) - \mathbb{E}_D[f_n(X)])^2] + \mathbb{E}_D[(\mathbb{E}_D[f_n(X)] - f^*(X))^2]] \\ & \quad + 2\mathbb{E}_D[(f_n(X) - \mathbb{E}_D[f_n(X)])(\mathbb{E}_D[f_n(X)] - f^*(X))] \\ &= \mathbb{E}_X[\mathbb{E}_D[(f_n(X) - \mathbb{E}_D[f_n(X)])^2] + (\mathbb{E}_D[f_n(X)] - f^*(X))^2] \\ &= \mathbb{E}_X[\text{Var}(f_n(X)) + (\text{Bias}(f_n(X)))^2], \end{aligned}$$

# Bias-Variance Decomposition

Finally, the expected risk over all training datasets is thus given by:

$$\mathbb{E}_D[R(f_n)] = \mathbb{E}_{Y,X}[(Y - f^*(X))^2] + \mathbb{E}_X[(\text{Bias } f_n(X))^2] + \mathbb{E}_X[\text{Var } f_n(X)],$$

which is written as **(Noise)-Bias-Variance-Decomposition**, i.e.,

**expected loss = noise + variance + squared bias.**

Trade-off between **bias** and **variance**  
 corresponds to  
 Trade-off between **overfitting** and **underfitting**.

# Bias-Variance Decomposition

Finally, the expected risk over all training datasets is thus given by:

$$\mathbb{E}_D[R(f_n)] = \mathbb{E}_{Y,X}[(Y - f^*(X))^2] + \mathbb{E}_X[(\text{Bias } f_n(X))^2] + \mathbb{E}_X[\text{Var } f_n(X)],$$

which is written as **(Noise)-Bias-Variance-Decomposition**, i.e.,

**expected loss = noise + variance + squared bias.**

Trade-off between **bias** and **variance**

corresponds to

Trade-off between **overfitting** and **underfitting**.

For a particular feature vector  $\mathbf{x}$ , then

- **Noise term:**  $\mathbb{E}[(Y - f^*(X))^2 | X = \mathbf{x}],$
- **Variance of  $f_n$ :**  $\text{Var } f_n(\mathbf{x}) = \mathbb{E}_D[(f_n(\mathbf{x}) - \mathbb{E}_D[f_n(\mathbf{x})])^2],$
- **Bias of  $f_n$ :**  $\text{Bias } f_n(\mathbf{x}) = \mathbb{E}_D[f_n(\mathbf{x})] - f^*(\mathbf{x}),$

where the latter two terms correspond to the variance and the bias of the  $f_n$  estimator at  $X = \mathbf{x}$ .



# Bias-Variance of Least Squares

## Bias-Variance-Decomposition for the Least-Squares estimator:

the bias and variance of  $f_n(\mathbf{x}) = \langle \mathbf{w}_n, \mathbf{x} \rangle$  (i.e., the LS prediction of the target value  $y$  for a feature vector  $\mathbf{x}$ ) can be written in terms of respectively the bias and the covariance of  $\mathbf{w}_n$ , i.e.,

$$\begin{aligned} \text{Bias } f_n(X = \mathbf{x}) &= \mathbb{E}_D[f_n(\mathbf{x})] - f^*(X) = \mathbb{E}_D[\langle \mathbf{w}_n, \mathbf{x} \rangle] - \langle \mathbf{w}^*, \mathbf{x} \rangle \\ &= \langle \mathbb{E}_D[\mathbf{w}_n] - \mathbf{w}^*, \mathbf{x} \rangle \\ &= \langle \text{Bias}(\mathbf{w}_n), \mathbf{x} \rangle, \end{aligned}$$

$$\begin{aligned} \text{Var } f_n(X = \mathbf{x}) &= \mathbb{E}_D[(f_n(\mathbf{x}) - \mathbb{E}_D[f_n(\mathbf{x})])^2] = \mathbb{E}_D[(\langle \mathbf{w}_n, \mathbf{x} \rangle - \langle \mathbb{E}_D[\mathbf{w}_n], \mathbf{x} \rangle)^2] \\ &= \mathbb{E}_D[\langle \mathbf{w}_n - \mathbb{E}_D[\mathbf{w}_n], \mathbf{x} \rangle^2] \\ &= \mathbb{E}_D\left[\sum_{i,j=1}^d (w_{ni} - \mathbb{E}_D[w_{ni}])x_i x_j (w_{nj} - \mathbb{E}_D[w_{nj}])\right] \\ &= \text{tr}(\mathbf{x}\mathbf{x}^T \text{Cov}(\mathbf{w}_n)) \end{aligned}$$

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem**
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 Appendix

## 00

## Theorem (Gauss-Markov theorem)

*Suppose that the data obeys the linear model*

$$Y = \langle \mathbf{w}, X \rangle + \epsilon,$$

with  $\mathbb{E}[\epsilon|X = \mathbf{x}] = 0$ ,  $\text{Var}[\epsilon|X = \mathbf{x}] = \sigma^2$  and errors at different points are uncorrelated (i.e., noise is homoscedastic). Then,

- the **least squares estimator (LSE)**  $\mathbf{w}_n = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y}$  is **unbiased**, and
- among all possible unbiased estimators of the LSE  $\mathbf{w}_n$  has the **smallest variance**.

# Observations

The Gauss-Markov-Theorem is only of limited practical use:

- Model assumption has to be true! In reality linearity assumption is not often fulfilled.
- If the model assumption is correct, then the least squares estimator is the best among all possible **unbiased** estimators.
- However, a slightly biased estimator (e.g. ridge regression or lasso, covered in next lecture) may present a much smaller variance, and thus better expected squared error.
- In other words, biased estimators (with low variance) may be preferred in practice.

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 Appendix

# Ridge Regression

**Ridge regression:** Adds to the LSR formulation an  $L_2$ -Regularization term.

## Definition (Ridge regression)

Given sample  $D_n = (X_i, Y_i)_{i=1}^n$ , **ridge regression** is formulated as:

$$D_n \mapsto f_n = \arg \min_{f \in \mathcal{F}} \frac{1}{n} \sum_{i=1}^n (Y_i - \langle \mathbf{w}, \Phi(X_i) \rangle)^2 + \lambda \|\mathbf{w}\|_2^2.$$

- Provides a regularized version of LSR and thus will be less prone to overfitting.
- Even if the solution of LSR is not unique, Ridge regression has a unique solution.

# Solution of ridge regression

For a given dataset  $D_n = (\mathbf{x}_i, y_i)_{i=1}^n$ , and a basis function  $\Phi$  si that the feature vectors are transformed and collected in a matrix  $\Phi \in \Re^{n \times d}$ . Then, the solution of ridge regression is given by:

$$\mathbf{w}_{n,\lambda} = (\Phi^T \Phi + \lambda \mathbb{1}_d)^{-1} \Phi^T \mathbf{Y}.$$

## Properties:

- The solution  $\mathbf{w}_{n,\lambda}$  exists and is unique.
- The regularizer  $\Omega(\mathbf{w}) = \|\mathbf{w}\|_2^2$  corresponds to a Gaussian prior (with zero mean and unit variance) for maximum a posteriori (MAP) estimation, i.e.,

$$p(\mathbf{w}) \propto e^{-\Omega(\mathbf{w})} = e^{-\|\mathbf{w}\|_2^2}.$$


# LSR vs. Ridge regression I

The solution of least squared regression:

$$\mathbf{w}_n = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{Y}.$$

The solution of ridge regression:

$$\mathbf{w}_{n,\lambda} = (\Phi^T \Phi + \lambda \mathbb{1}_d)^{-1} \Phi^T \mathbf{Y}.$$



**Figure:** Linear least squares regression versus linear ridge regression. The regression function is linear. (Image from Prof. Hein)




# LSR vs. Ridge regression II

The solution of least squared regression:

$$\mathbf{w}_n = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{Y}.$$

The solution of ridge regression:

$$\mathbf{w}_{n,\lambda} = (\Phi^T \Phi + \lambda \mathbb{I}_d)^{-1} \Phi^T \mathbf{Y}.$$



**Figure:** Comparison of least squares and ridge regression using a set of periodic basis functions. (Image from Prof. Hein)


# Geometric interpretation I

**Linear LSR** uses SVD of  $X$  (i.e.,  $\mathbf{X} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$ ), where  $\text{rank}(\mathbf{\Sigma})=r$ ,

$$\mathbf{X}\mathbf{w}_n = \mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T\mathbf{V}^T(\mathbf{\Sigma}^+)^2\mathbf{V}^T\mathbf{V}\mathbf{\Sigma}^T\mathbf{U}^T\mathbf{Y} = \sum_{i=1}^r \mathbf{u}_i \langle \mathbf{u}_i, \mathbf{Y} \rangle,$$

i.e., the outputs are projected on the basis spanned by  $\mathbf{U}$ . Above,

$\mathbf{\Sigma}^+ \in \mathbb{R}^{n \times d}$  is defined as  $\Sigma_{ij}^+ = \begin{cases} 1/\sigma_i & \text{if } i=j \text{ and } i \leq r, \\ 0 & \text{otherwise} \end{cases}$



**FIGURE 3.2.** The  $N$ -dimensional geometry of least squares regression with two predictors. The outcome vector  $\mathbf{y}$  is orthogonally projected onto the hyperplane spanned by the input vectors  $\mathbf{x}_1$  and  $\mathbf{x}_2$ . The projection  $\hat{\mathbf{y}}$  represents the vector of the least squares predictions

**Figure:** Image from ESL.

(Refer to 9 Appendix for further details on SVD.)

# Geometric interpretation II

**Linear LSR** uses SVD of  $X$  (i.e.,  $\mathbf{X} = \mathbf{U}\Sigma\mathbf{V}^T$ ), where  $\text{rank}(\Sigma)=r$ ,

$$\mathbf{X}\mathbf{w}_n = \mathbf{X}(\mathbf{X}^T\mathbf{X})^{-1}\mathbf{X}^T\mathbf{Y} = \mathbf{U}\Sigma\mathbf{V}^T\mathbf{V}^T(\Sigma^+)^2\mathbf{V}^T\mathbf{V}\Sigma^T\mathbf{U}^T\mathbf{Y} = \sum_{i=1}^r \mathbf{u}_i \langle \mathbf{u}_i, \mathbf{Y} \rangle,$$

i.e., the outputs are projected on the basis spanned by  $\mathbf{U}$ . Above,

$\Sigma^+ \in \mathbb{R}^{n \times d}$  is defined as  $\Sigma_{ij}^+ = \begin{cases} 1/\sigma_i & \text{if } i=j \text{ and } i \leq r, \\ 0 & \text{otherwise} \end{cases}$

**Ridge regression:**

$$\mathbf{X}\mathbf{w}_{n,\lambda} = \mathbf{X}(\mathbf{X}^T\mathbf{X} + \lambda\mathbb{I}_d)^{-1}\mathbf{X}^T\mathbf{Y} = \sum_{i=1}^r \mathbf{u}_i \frac{\sigma_i^2}{\sigma_i^2 + \lambda} \langle \mathbf{u}_i, \mathbf{Y} \rangle,$$

i.e., the output  $\mathbf{Y}$  is also projected on the basis spanned by  $\mathbf{U}$ , but here the smaller the singular value  $\sigma_i$  (compared to  $\lambda$ ) the larger the **shrinkage** in this direction, or in other words, the smaller is the influence of this direction.

(Refer to 9 Appendix for further details on SVD.)

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso**
- 8 Summary
- 9 Appendix

# Lasso

**Lasso** (least absolute shrinkage and selection operator) corresponds to LSR with  $L_1$ -Regularization.

## Definition (Lasso)

Given a training sample  $D_n = (X_i, Y_i)_{i=1}^n$  with  $X_i \in \mathcal{X}$  and  $Y_i \in \mathbb{R}$  and the function space  $\mathcal{F} = \{\sum_{j=1}^d w_j \phi_j(\mathbf{x}) \mid \mathbf{w} \in \mathbb{R}^d\}$  we define **the lasso** as the mapping  $\mathcal{A} : D_n \rightarrow \mathcal{F}$  with


$$D_n \mapsto \mathbf{w}_n = \arg \min_{\mathbf{w} \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n (Y_i - \langle \mathbf{w}, \Phi(X_i) \rangle)^2 + \lambda \|\mathbf{w}\|_1.$$

*Observation:* The regularizer  $\Omega(\mathbf{w}) = \|\mathbf{w}\|_1$  corresponds to a Laplace prior for maximum a posteriori (MAP) estimation.

# $L_1$ vs other norms


- A closed form **solution** as for Lasso is generally not available. However, as the objective function to be minimized is convex with respect to  $\mathbf{w}$ , one may rely on existing and efficient convex optimization techniques to compute the solution to Lasso.
- $L_1$ -norm induces **sparsity** (a lot of elements of  $\mathbf{w}_n$  are zero). Sparsity is good: less storage, faster evaluation  $f(\mathbf{x}) = \langle \mathbf{w}, \mathbf{x} \rangle$ , feature selection.
- Why?  $L_1$ -norm is the norm which is “closest” to the “zero norm” ( $\|\mathbf{w}\|_0 = \sum_{i=1}^D \mathbb{1}_{w_i \neq 0}$ ). The “zero norm” enforces directly sparsity.
- $L_2$ -norm  $\|\mathbf{w}\|_2$  penalizes large weights heavily  $\Rightarrow$  preference for small weights in all directions (regularizer is **isotropic**).  
 $L_1$ -norm  $\|\mathbf{w}\|_1$  penalizes large and small weights “equally”  $\Rightarrow$  produces often large weights in few directions.

## Comparison: lasso and ridge regression I



**Figure:** **Left:** Perturbed training data and regression function in black, we show the solution of ridge regression in blue and of Lasso in red for  $\lambda = 1$ , **Right:** Behavior of training and test error and number of non-zero components of the weight vector as a function of the regularization parameter  $\lambda$ . (Images from Prof. Hein)

## Comparison: lasso and ridge regression II



**Figure: Left:** The ridge weights (coefficients) present smaller absolute values than for LSR, but still different that zero. **Right:** In contrast, Lasso triggers a subset of weights directly to zero. (Images from eML)



## Comparison: lasso and ridge regression II

**Figure 3.4** Plot of the contours of the unregularized error function (blue) along with the constraint region (3.30) for the quadratic regularizer  $q = 2$  on the left and the lasso regularizer  $q = 1$  on the right, in which the optimum value for the parameter vector  $\mathbf{w}$  is denoted by  $\mathbf{w}^*$ . The lasso gives a sparse solution in which  $w_1^* = 0$ .





Figure: Image from Bishop.

# Regularization functionals

**Other regularization functionals:**  $\Omega(\mathbf{w}) = \sum_{i=1}^n |w_i|^p = \|\mathbf{w}\|_p^p$ .  
 $\Rightarrow$   $L_2$ -norm is the only **isotropic** norm in the family of  $p$ -norms!



**Figure:** The level set  $\|\mathbf{w}\|_p = 1$  of the  $p$ -norms. Note that the  $\|\cdot\|_p$  is only a norm for  $p \geq 1$ , in which case the unit-ball is a convex set. Clearly for  $p = 0.5$  the “unit-ball” is not convex.

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 **Summary**
- 9 Appendix

# Summary

**Linear regression** (and as we will see linear classification too):

- Easy interpretation: feature has a high influence if it has a large weight.
- Linear methods: have possibly high bias but low variance  $\Rightarrow$  can be fit already with only a few training points.
- Often competitive with non-linear methods in high dimensions.
- Using transformations of the input features (**basis functions**) one can easily generate non-linear functions in the input space. Linear methods are *linear* in the parameters, but not necessarily linear in the original input features.

# Summary II

## Regularized Linear Regression:

- Adding a regularization term, e.g., L2-norm, to LSR formulation allows us to control for overfitting (i.e., for the bias-variance trade-off). It also makes the solution to the regression problem unique, even when the LSR solution is not.
- Lasso induces sparsity in the regression function weights (i.e., it forces the value of a subset of weights to zero). Thus, Lasso leads to stronger regularization than Ridge regression. However, in contrast to Ridge regression, Lasso does not have close-form solution (efficient convex optimization techniques exist).
- Ridge regression and Lasso can be seen a MAP estimation of the regression weights with, respectively, Gaussian and Laplace priors.

# Outline

- 1 Bibliography
- 2 Least squares regression
- 3 Basis functions
- 4 Bias-variance
- 5 Gauss-Markov-Theorem
- 6 Ridge Regression
- 7 Lasso
- 8 Summary
- 9 **Appendix**

# The pseudo-inverse

If  $\mathbf{X}$  has rank  $d$ , then  $(\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T$  is the Moore-Penrose **pseudo inverse** of  $\mathbf{X}$ .

### Definition (Pseudo-inverse)

Let  $\mathbf{A} \in \mathbb{R}^{m \times n}$  with rank  $r \leq \min\{m, n\}$ . Then the **pseudo-inverse**  $\mathbf{A}^+$  of  $\mathbf{A}$  is defined as

$$\mathbf{A}^+ = \arg \min_{\mathbf{B} \in \mathbb{R}^{n \times m}} \|\mathbf{AB} - \mathbb{1}_m\|_F^2,$$

where  $\|\cdot\|_F$  is the **Frobenius norm** ( $\|\mathbf{A}\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n A_{ij}^2}$ ) and  $\mathbb{1}_m$  the identity matrix in  $\mathbb{R}^m$ .

Let  $\mathbf{A}$  be an invertible square matrix, then

$$(\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T = \mathbf{A}^{-1} (\mathbf{A}^T)^{-1} \mathbf{A}^T = \mathbf{A}^{-1}.$$

## SVD

The **singular value decomposition** of  $\mathbf{A} \in \mathbb{R}^{m \times n}$ ,

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T,$$

- $\mathbf{U}$  is an orthogonal matrix  $\mathbf{U} = (\mathbf{u}_1, \dots, \mathbf{u}_m) \in \mathbb{R}^{m \times m}$ , that is  $\mathbf{U}^T \mathbf{U} = \mathbf{1}_m$ ,
- $\mathbf{V}$  is an orthogonal matrix  $\mathbf{V} = (\mathbf{v}_1, \dots, \mathbf{v}_n) \in \mathbb{R}^{n \times n}$ , that is  $\mathbf{V}^T \mathbf{V} = \mathbf{1}_n$ ,
- $\mathbf{\Sigma} \in \mathbb{R}^{m \times n}$  with  $\Sigma_{ij} = \begin{cases} \sigma_i & \text{if } i = j \text{ and } i \leq r, \\ 0 & \text{otherwise} \end{cases}$ .

The  $\sigma_i > 0$ ,  $i = 1, \dots, r$  are the **singular values** of  $\mathbf{A}$ .

The **pseudo inverse** of a matrix  $\mathbf{A}$  is then given by

$$\mathbf{A}^+ = \mathbf{A}^{-1}(\mathbf{A}^T)^{-1}\mathbf{A}^T = \mathbf{V}\mathbf{\Sigma}^+\mathbf{U}^T,$$

where  $\mathbf{\Sigma}^+ \in \mathbb{R}^{n \times m}$  is defined as  $\Sigma_{ij}^+ = \begin{cases} 1/\sigma_i & \text{if } i = j \text{ and } i \leq r, \\ 0 & \text{otherwise} \end{cases}$ .



# Proof.

The **pseudo inverse**  $A^+$  is then given by

$$\mathbf{A}^+ = \mathbf{V}\mathbf{\Sigma}^+\mathbf{U}^T,$$

where  $\mathbf{\Sigma}^+ \in \mathbb{R}^{n \times m}$  is given by  $\Sigma_{ij}^+ = \begin{cases} 1/\sigma_i & \text{if } i = j \text{ and } i \leq r, \\ 0 & \text{otherwise} \end{cases}$ .

Let  $\mathbf{A} \in \mathbb{R}^{n \times m}$ . Given that  $m \leq n$  and  $\text{rank}(\mathbf{A}) = m$ , one can write the pseudo inverse  $\mathbf{A}^+$  as  $\mathbf{A}^+ = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T$ ,

$$\begin{aligned} (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T &= (\mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T \mathbf{U} \mathbf{\Sigma} \mathbf{V}^T)^{-1} \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T = (\mathbf{V} \mathbf{\Sigma}^T \mathbf{\Sigma} \mathbf{V}^T)^{-1} \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T \\ &= \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{V}^T \mathbf{V} \mathbf{\Sigma}^T \mathbf{U}^T = \mathbf{V} (\mathbf{\Sigma}^T \mathbf{\Sigma})^{-1} \mathbf{\Sigma}^T \mathbf{U}^T = \mathbf{V} \mathbf{\Sigma}^+ \mathbf{U}^T. \end{aligned}$$